

# VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



**LAB REPORT on**

## **BIG DATA ANALYTICS (23CS6PEBDA)**

*Submitted by*

**AMRUTHA MURALIDHAR (1BM21CS257)**

*in partial fulfilment for the award of the degree of*  
**BACHELOR OF ENGINEERING**

*in*

**COMPUTER SCIENCE AND ENGINEERING**



**B.M.S. COLLEGE OF ENGINEERING BENGALURU-560019 May-2022 to  
July-2022**

(Autonomous Institution under VTU)

**B. M. S. College of Engineering,  
Bull Temple Road, Bangalore 560019**

(Affiliated To Visvesvaraya Technological University, Belgaum)

**Department of Computer Science and Engineering**



**CERTIFICATE**

This is to certify that the Lab work entitled “**BIG DATA ANALYTICS** ” carried out by **Amrutha Muralidhar (1BM21CS257)**, who is a bonafide student of **B. M. S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum during the year 2023. The Lab report has been approved as it satisfies the academic requirements in respect of Big data analytics - (23CS6PEBDA) work prescribed for the said degree.

Lab-In charge  
Rekha G S  
Assistant Professor  
Department of CSE  
BMSCE, Bengaluru

## Index

<b>Sl. No</b>	<b>Experiment Title</b>	<b>Page No.</b>
1	Cassandra Lab Program 1: - Employee Database	5
2	Cassandra Lab Program 2: - Library Database	9
3	MongoDB- CRUD Demonstration	15
4	Hadoop Installation	22
5	Hadoop Commands	23
6	Hadoop Programs: Word Count	26
7	Hadoop Programs: Top N	32
8	Hadoop Programs: Average Temperature	37

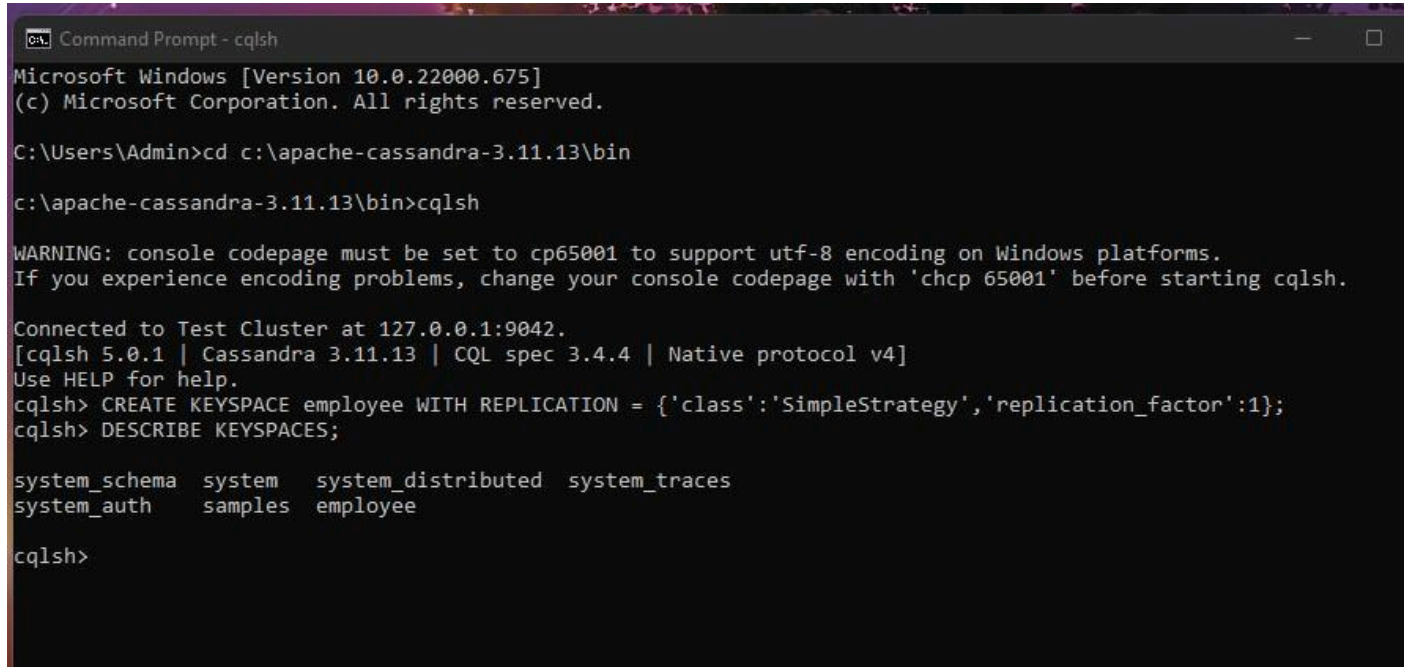
## Course Outcome

CO1	Apply the concept of NoSQL, Hadoop or Spark for a given task
CO2	Analyze the Big Data and obtain insight using data analytics mechanisms.
CO3	Design and implement Big data applications by applying NoSQL, Hadoop or Spark

## Cassandra Lab Program 1: -

Perform the following DB operations using Cassandra.

1. Create a key space by name Employee



```
Command Prompt - cqlsh
Microsoft Windows [Version 10.0.22000.675]
(c) Microsoft Corporation. All rights reserved.

C:\Users\Admin>cd c:\apache-cassandra-3.11.13\bin

c:\apache-cassandra-3.11.13\bin>cqlsh

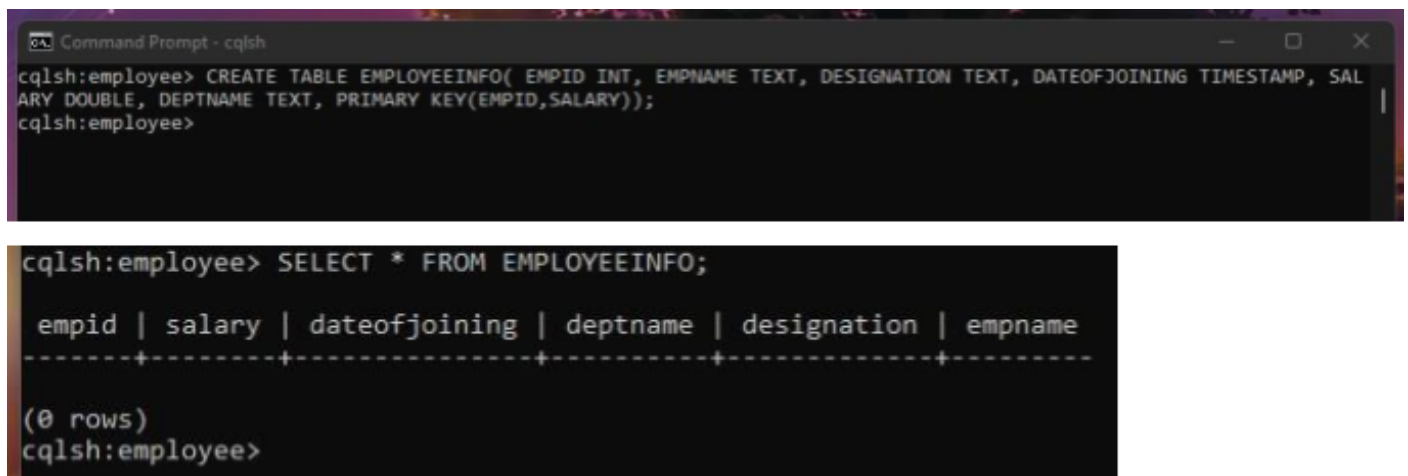
WARNING: console codepage must be set to cp65001 to support utf-8 encoding on Windows platforms.
If you experience encoding problems, change your console codepage with 'chcp 65001' before starting cqlsh.

Connected to Test Cluster at 127.0.0.1:9042.
[cqlsh 5.0.1 | Cassandra 3.11.13 | CQL spec 3.4.4 | Native protocol v4]
Use HELP for help.
cqlsh> CREATE KEYSPACE employee WITH REPLICATION = {'class':'SimpleStrategy','replication_factor':1};
cqlsh> DESCRIBE KEYSPACES;

system_schema  system  system_distributed  system_traces
system_auth    samples  employee

cqlsh>
```

2. Create a column family by name Employee-Info with attributes Emp\_Id Primary Key, Emp\_Name, Designation, Date\_of\_Joining, Salary, Dept\_Name



```
Command Prompt - cqlsh

cqlsh:employee> CREATE TABLE EMPLOYEEINFO( EMPID INT, EMPNAME TEXT, DESIGNATION TEXT, DATEOFJOINING TIMESTAMP, SALARY DOUBLE, DEPTNAME TEXT, PRIMARY KEY(EMPID,SALARY));
cqlsh:employee>

cqlsh:employee> SELECT * FROM EMPLOYEEINFO;

 empid | salary | dateofjoining | deptname | designation | empname
-----+-----+-----+-----+-----+-----
(0 rows)
cqlsh:employee>
```

3. Insert the values into the table in batch

```

cqlsh:employee> BEGIN BATCH
... INSERT INTO EMPLOYEEINFO (EMPID, EMPNAME, DESIGNATION, DATEOFJOINING, SALARY, DEPTNAME)
... VALUES(1,'LOKESH','ASSISTANT MANAGER', '2005-04-6', 50000, 'MARKETING')
... INSERT INTO EMPLOYEEINFO (EMPID, EMPNAME, DESIGNATION, DATEOFJOINING, SALARY, DEPTNAME)
... VALUES(2,'DHEERAJ','ASSISTANT MANAGER', '2013-11-10', 30000, 'LOGISTICS')
... INSERT INTO EMPLOYEEINFO (EMPID, EMPNAME, DESIGNATION, DATEOFJOINING, SALARY, DEPTNAME)
... VALUES(3,'CHIRAG','ASSISTANT MANAGER', '2011-07-1', 115000, 'SALES')
... INSERT INTO EMPLOYEEINFO (EMPID, EMPNAME, DESIGNATION, DATEOFJOINING, SALARY, DEPTNAME)
... VALUES(4,'DHANUSH','ASSISTANT MANAGER', '2010-04-26', 75000, 'MARKETING')
... INSERT INTO EMPLOYEEINFO (EMPID, EMPNAME, DESIGNATION, DATEOFJOINING, SALARY, DEPTNAME)
... VALUES(5,'ESHA','ASSISTANT MANAGER', '2010-04-26', 85000, 'TECHNICAL')
... INSERT INTO EMPLOYEEINFO (EMPID, EMPNAME, DESIGNATION, DATEOFJOINING, SALARY, DEPTNAME)
... VALUES(6,'FARHAN','MANAGER', '2010-04-26', 95000, 'TECHNICAL')
... INSERT INTO EMPLOYEEINFO (EMPID, EMPNAME, DESIGNATION, DATEOFJOINING, SALARY, DEPTNAME)
... VALUES(7,'JIMMY','MANAGER', '2010-04-26', 95000, 'PR')
... INSERT INTO EMPLOYEEINFO (EMPID, EMPNAME, DESIGNATION, DATEOFJOINING, SALARY, DEPTNAME)
... VALUES(121,'HARRY','REGIONAL MANAGER', '2010-04-26', 99000, 'MANAGEMENT')
... APPLY BATCH;

cqlsh:employee> SELECT * FROM EMPLOYEEINFO;

empid | salary | dateofjoining | deptname | designation | empname
-----+-----+-----+-----+-----+-----
5 | 85000 | 2010-04-25 18:30:00.000000+0000 | TECHNICAL | ASSISTANT MANAGER | ESHA
1 | 50000 | 2005-04-05 18:30:00.000000+0000 | MARKETING | ASSISTANT MANAGER | LOKESH
2 | 30000 | 2013-11-09 18:30:00.000000+0000 | LOGISTICS | ASSISTANT MANAGER | DHEERAJ
4 | 75000 | 2010-04-25 18:30:00.000000+0000 | MARKETING | ASSISTANT MANAGER | DHANUSH
121 | 99000 | 2010-04-25 18:30:00.000000+0000 | MANAGEMENT | REGIONAL MANAGER | HARRY
7 | 95000 | 2010-04-25 18:30:00.000000+0000 | PR | MANAGER | JIMMY
6 | 95000 | 2010-04-25 18:30:00.000000+0000 | TECHNICAL | MANAGER | FARHAN
3 | 1.15e+05 | 2011-06-30 18:30:00.000000+0000 | SALES | ASSISTANT MANAGER | CHIRAG

(8 rows)
cqlsh:employee>

```

#### 4. Update Employee name and Department of Emp-Id 121

```

cqlsh:employee> UPDATE EMPLOYEEINFO SET EMPNAME='HARRY', DEPTNAME='MANAGEMENT' WHERE EMPID=121 AND SALARY=99000;
cqlsh:employee> SELECT * FROM EMPLOYEEINFO;

empid | salary | dateofjoining | deptname | designation | empname
-----+-----+-----+-----+-----+-----
5 | 85000 | 2010-04-25 18:30:00.000000+0000 | TECHNICAL | ASSISTANT MANAGER | ESHA
1 | 50000 | 2005-04-05 18:30:00.000000+0000 | MARKETING | ASSISTANT MANAGER | LOKESH
2 | 30000 | 2013-11-09 18:30:00.000000+0000 | LOGISTICS | ASSISTANT MANAGER | DHEERAJ
4 | 75000 | 2010-04-25 18:30:00.000000+0000 | MARKETING | ASSISTANT MANAGER | DHANUSH
121 | 99000 | 2010-04-25 18:30:00.000000+0000 | MANAGEMENT | REGIONAL MANAGER | HARRY
7 | 95000 | 2010-04-25 18:30:00.000000+0000 | PR | MANAGER | JIMMY
6 | 95000 | 2010-04-25 18:30:00.000000+0000 | TECHNICAL | MANAGER | FARHAN
3 | 1.15e+05 | 2011-06-30 18:30:00.000000+0000 | SALES | ASSISTANT MANAGER | CHIRAG

(8 rows)
cqlsh:employee>

```

#### 5. Sort the details of Employee records based on salary (Note:- cql>PAGING OFF)



```
cqlsh:employee> select * from EMPLOYEEINFO where empid IN(1,2,3,4,5,6,7) ORDER BY salary DESC allow filtering;
```

empid	salary	dateofjoining	deptname	designation	empname
3	1.15e+05	2011-06-30 18:30:00.000000+0000	SALES	ASSISTANT MANAGER	CHIRAG
6	95000	2010-04-25 18:30:00.000000+0000	TECHNICAL	MANAGER	FARHAN
7	95000	2010-04-25 18:30:00.000000+0000	PR	MANAGER	JIMMY
5	85000	2010-04-25 18:30:00.000000+0000	TECHNICAL	ASSISTANT MANAGER	ESHA
4	75000	2010-04-25 18:30:00.000000+0000	MARKETING	ASSISTANT MANAGER	DHANUSH
1	50000	2005-04-05 18:30:00.000000+0000	MARKETING	ASSISTANT MANAGER	LOKESH
2	30000	2013-11-09 18:30:00.000000+0000	LOGISTICS	ASSISTANT MANAGER	DHEERAJ

```
(7 rows)
cqlsh:employee>
```

6. Alter the schema of the table Employee\_Info to add a column Projects which stores a set of Projects done by the corresponding Employee.

```
(7 rows)
cqlsh:employee> ALTER TABLE EMPLOYEEINFO ADD PROJECTS LIST<TEXT>;
cqlsh:employee> SELECT * FROM EMPLOYEEINFO;
```

empid	salary	dateofjoining	deptname	designation	empname	projects
5	85000	2010-04-25 18:30:00.000000+0000	TECHNICAL	ASSISTANT MANAGER	ESHA	null
1	50000	2005-04-05 18:30:00.000000+0000	MARKETING	ASSISTANT MANAGER	LOKESH	null
2	30000	2013-11-09 18:30:00.000000+0000	LOGISTICS	ASSISTANT MANAGER	DHEERAJ	null
4	75000	2010-04-25 18:30:00.000000+0000	MARKETING	ASSISTANT MANAGER	DHANUSH	null
121	99000	2010-04-25 18:30:00.000000+0000	MANAGEMENT	REGIONAL MANAGER	HARRY	null
7	95000	2010-04-25 18:30:00.000000+0000	PR	MANAGER	JIMMY	null
6	95000	2010-04-25 18:30:00.000000+0000	TECHNICAL	MANAGER	FARHAN	null
3	1.15e+05	2011-06-30 18:30:00.000000+0000	SALES	ASSISTANT MANAGER	CHIRAG	null

```
(8 rows)
cqlsh:employee> _
```

7. Update the altered table to add project names.

```
Command Prompt - cqlsh
cqlsh:employee> UPDATE EMPLOYEEINFO SET PROJECTS=['FACEBOOK','SNAPCHAT'] WHERE EMPID=1 AND SALARY=50000;
cqlsh:employee> UPDATE EMPLOYEEINFO SET PROJECTS=['FACEBOOK','SNAPCHAT'] WHERE EMPID=7 AND SALARY=95000;
cqlsh:employee> UPDATE EMPLOYEEINFO SET PROJECTS=['PINTEREST','INSTAGRAM'] WHERE EMPID=121 AND SALARY=99000;
cqlsh:employee> UPDATE EMPLOYEEINFO SET PROJECTS=['PINTEREST','INSTAGRAM'] WHERE EMPID=4 AND SALARY=75000;
cqlsh:employee> UPDATE EMPLOYEEINFO SET PROJECTS=['YOUTUBE','SPOTIFY'] WHERE EMPID=2 AND SALARY=30000;
cqlsh:employee> UPDATE EMPLOYEEINFO SET PROJECTS=['YOUTUBE','SPOTIFY'] WHERE EMPID=3 AND SALARY=115000;
cqlsh:employee> UPDATE EMPLOYEEINFO SET PROJECTS=['YOUTUBE','SPOTIFY'] WHERE EMPID=6 AND SALARY=95000;
cqlsh:employee> UPDATE EMPLOYEEINFO SET PROJECTS=['YOUTUBE','SPOTIFY'] WHERE EMPID=5 AND SALARY=85000;
cqlsh:employee> SELECT * FROM EMPLOYEEINFO;
```

empid	salary	dateofjoining	deptname	designation	empname	projects
5	85000	2010-04-25 18:30:00.000000+0000	TECHNICAL	ASSISTANT MANAGER	ESHA	['YOUTUBE', 'SPOTIFY']
1	50000	2005-04-05 18:30:00.000000+0000	MARKETING	ASSISTANT MANAGER	LOKESH	['FACEBOOK', 'SNAPCHAT']
2	30000	2013-11-09 18:30:00.000000+0000	LOGISTICS	ASSISTANT MANAGER	DHEERAJ	['YOUTUBE', 'SPOTIFY']
4	75000	2010-04-25 18:30:00.000000+0000	MARKETING	ASSISTANT MANAGER	DHANUSH	['PINTEREST', 'INSTAGRAM']
121	99000	2010-04-25 18:30:00.000000+0000	MANAGEMENT	REGIONAL MANAGER	HARRY	['PINTEREST', 'INSTAGRAM']
7	95000	2010-04-25 18:30:00.000000+0000	PR	MANAGER	JIMMY	['FACEBOOK', 'SNAPCHAT']
6	95000	2010-04-25 18:30:00.000000+0000	TECHNICAL	MANAGER	FARHAN	['YOUTUBE', 'SPOTIFY']
3	1.15e+05	2011-06-30 18:30:00.000000+0000	SALES	ASSISTANT MANAGER	CHIRAG	['YOUTUBE', 'SPOTIFY']

```
(8 rows)
cqlsh:employee>
```

8. Create a TTL of 15 seconds to display the values of Employees.

//BEFORE 15 seconds

```
Command Prompt - cqlsh
cqlsh:employee> update EMPLOYEEINFO USING TTL 15 SET EMPNAME='LOKESH' where empid=1 AND salary=50000;
cqlsh:employee> SELECT * FROM EMPLOYEEINFO;
```

empid	salary	dateofjoining	deptname	designation	empname	projects
5	85000	2010-04-25 18:30:00.000000+0000	TECHNICAL	ASSISTANT MANAGER	ESHA	['YOUTUBE', 'SPOTIFY']
1	50000	2005-04-05 18:30:00.000000+0000	MARKETING	ASSISTANT MANAGER	LOKESH	['FACEBOOK', 'SNAPCHAT']
2	30000	2013-11-09 18:30:00.000000+0000	LOGISTICS	ASSISTANT MANAGER	DHEERAJ	['YOUTUBE', 'SPOTIFY']
4	75000	2010-04-25 18:30:00.000000+0000	MARKETING	ASSISTANT MANAGER	DHANUSH	['PINTEREST', 'INSTAGRAM']
121	99000	2010-04-25 18:30:00.000000+0000	MANAGEMENT	REGIONAL MANAGER	HARRY	['PINTEREST', 'INSTAGRAM']
7	95000	2010-04-25 18:30:00.000000+0000	PR	MANAGER	JIMMY	['FACEBOOK', 'SNAPCHAT']
6	95000	2010-04-25 18:30:00.000000+0000	TECHNICAL	MANAGER	FARHAN	['YOUTUBE', 'SPOTIFY']
3	1.15e+05	2011-06-30 18:30:00.000000+0000	SALES	ASSISTANT MANAGER	CHIRAG	['YOUTUBE', 'SPOTIFY']

(8 rows)  
cqlsh:employee>

## Cassandra Lab Program 2: -

Perform the following DB operations using Cassandra.

1. Create a key space by name Library

```
Command Prompt - CQLSH
cqlsh> create keyspace library with replication = {
... 'class':'SimpleStrategy', 'replication_factor':1
... };
cqlsh> describe keyspaces

system_schema  system  samples  employee
system_auth    library system_distributed system_traces

cqlsh> USE library;
cqlsh:library> _
```

2. Create a column family by name Library-Info with attributes Stud\_Id Primary Key,

Counter\_value of type Counter,

Stud\_Name, Book-Name, Book-Id, Date\_of\_issue

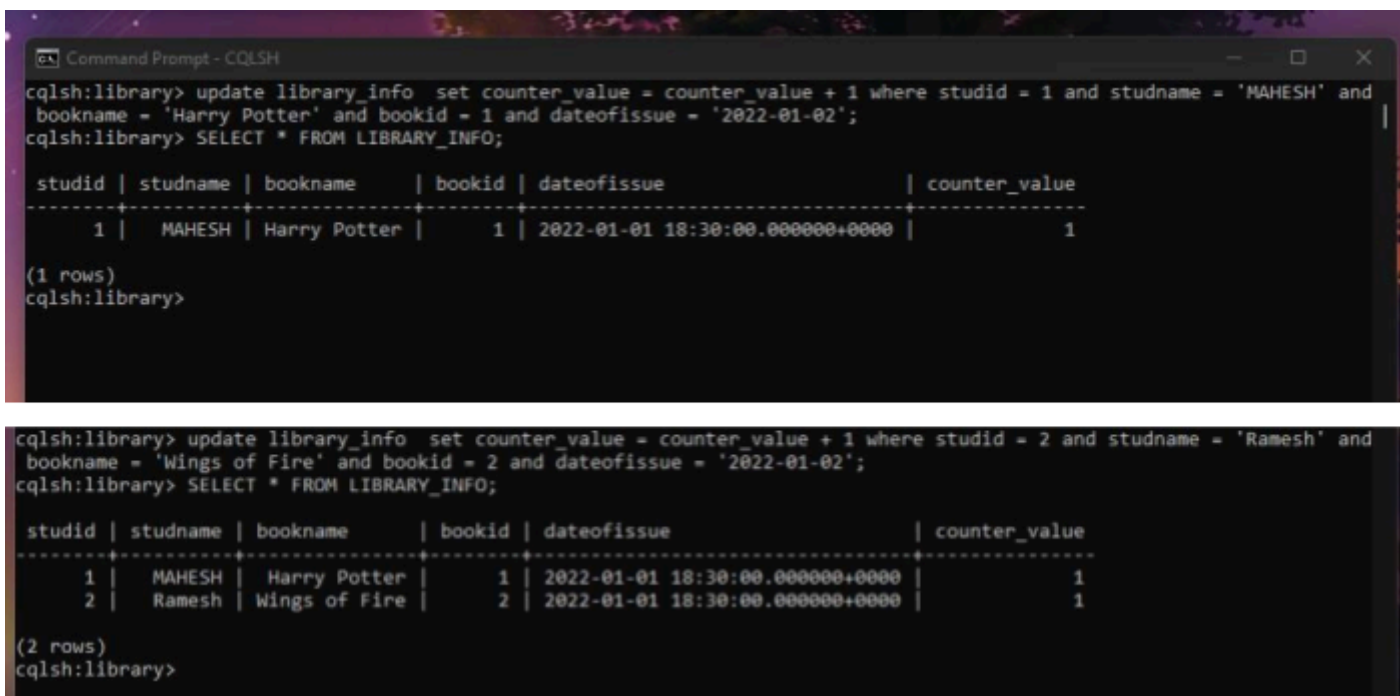
```

cqlsh> USE library;
cqlsh:library> CREATE TABLE LIBRARY_INFO( STUDID INT PRIMARY KEY, STUDNAME TEXT, BOOKNAME TEXT, DATEOFISSUE TIMESTAMP,
COUNTER_VALUE COUNTER);
InvalidRequest: Error from server: code=2200 [Invalid query] message="Cannot mix counter and non counter columns in the
same table"
cqlsh:library> CREATE TABLE LIBRARY_INFO( STUDID INT, STUDNAME TEXT, BOOKNAME TEXT, BOOKID INT, DATEOFISSUE TIMESTAMP,
COUNTER_VALUE COUNTER, PRIMARY KEY(STUDID, STUDNAME, BOOKNAME, BOOKID, DATEOFISSUE));
cqlsh:library> SELECT * FROM LIBRARYINFO;
InvalidRequest: Error from server: code=2200 [Invalid query] message="unconfigured table libraryinfo"
cqlsh:library> SELECT * FROM LIBRARY_INFO;

studid | studname | bookname | bookid | dateofissue | counter_value
-----+-----+-----+-----+-----+-----
(0 rows)
cqlsh:library>

```

3. Insert the values into the table in batch



```

Command Prompt - CQLSH
cqlsh:library> update library_info set counter_value = counter_value + 1 where studid = 1 and studname = 'MAHESH' and
bookname = 'Harry Potter' and bookid = 1 and dateofissue = '2022-01-02';
cqlsh:library> SELECT * FROM LIBRARY_INFO;

studid | studname | bookname | bookid | dateofissue | counter_value
-----+-----+-----+-----+-----+-----
1 | MAHESH | Harry Potter | 1 | 2022-01-01 18:30:00.000000+0000 | 1

(1 rows)
cqlsh:library>

cqlsh:library> update library_info set counter_value = counter_value + 1 where studid = 2 and studname = 'Ramesh' and
bookname = 'Wings of Fire' and bookid = 2 and dateofissue = '2022-01-02';
cqlsh:library> SELECT * FROM LIBRARY_INFO;

studid | studname | bookname | bookid | dateofissue | counter_value
-----+-----+-----+-----+-----+-----
1 | MAHESH | Harry Potter | 1 | 2022-01-01 18:30:00.000000+0000 | 1
2 | Ramesh | Wings of Fire | 2 | 2022-01-01 18:30:00.000000+0000 | 1

(2 rows)
cqlsh:library>

```

4. Display the details of the table created and increase the value of the counter



```
cqlsh:library> update library_info set counter_value = counter_value + 1 where studid = 112 and studname = 'Rajesh' and bookname = 'BDA' and bookid = 3 and dateofissue = '2022-01-02';
cqlsh:library> SELECT * FROM LIBRARY_INFO;
```

studid	studname	bookname	bookid	dateofissue	counter_value
1	MAHESH	Harry Potter	1	2022-01-01 18:30:00.000000+0000	1
2	Ramesh	Wings of Fire	2	2022-01-01 18:30:00.000000+0000	1
112	Rajesh	BDA	3	2022-01-01 18:30:00.000000+0000	1

(3 rows)  
cqlsh:library>

```
(3 rows)
cqlsh:library> update library_info set counter_value = counter_value + 1 where studid = 112 and studname = 'Rajesh' and bookname = 'BDA' and bookid = 3 and dateofissue = '2022-01-02';
cqlsh:library> SELECT * FROM LIBRARY_INFO;
```

studid	studname	bookname	bookid	dateofissue	counter_value
1	MAHESH	Harry Potter	1	2022-01-01 18:30:00.000000+0000	1
2	Ramesh	Wings of Fire	2	2022-01-01 18:30:00.000000+0000	1
112	Rajesh	BDA	3	2022-01-01 18:30:00.000000+0000	2

(3 rows)  
cqlsh:library>

studid	studname	bookname	bookid	dateofissue	counter_value
113	Ranjith	rpa	4	2022-01-01 18:30:00.000000+0000	1
1	MAHESH	Harry Potter	1	2022-01-01 18:30:00.000000+0000	1
2	Ramesh	Wings of Fire	2	2022-01-01 18:30:00.000000+0000	1
112	Rajesh	BDA	3	2022-01-01 18:30:00.000000+0000	3

(4 rows)

- Write a query to show that a student with id 112 has taken a book “BDA” 3 times.

```
Command Prompt - CQLSH
cqlsh:library> select * from library_info where studid = 112;
```

studid	studname	bookname	bookid	dateofissue	counter_value
112	Rajesh	BDA	3	2022-01-01 18:30:00.000000+0000	3

(1 rows)  
cqlsh:library>

- Export the created column to a csv file


```
cqlsh:library> copy library_info (studid, studname, bookname, bookid, dateofissuse, counter_value) to 'C:\Users\Admin\OneDrive\Desktop\BDA Lab\data.csv';
Using 7 child processes

Starting copy of library.library_info with columns [studid, studname, bookname, bookid, dateofissuse, counter_value].
Processed: 4 rows; Rate:      2 rows/s; Avg. rate:      1 rows/s
4 rows exported to 1 files in 3.004 seconds.
cqlsh:library> _
```




Clipboard

Font

Alignment

 POSSIBLE DATA LOSS Some features might be lost if you save this workbook in the comma-delimited

A1

113

	A	B	C	D	E	F	G	H	I
1	113	Ranjith	rpa		4 2022-01-01	1			
2	2	Ramesh	Wings of F		2 2022-01-01	1			
3	112	Rajesh	BDA		3 2022-01-01	3			
4	1	MAHESH	Harry Pott		1 2022-01-01	1			
5									
6									
7									

## 7. Import a given csv dataset from local file system into Cassandra column family

```
cqlsh:library> copy library_info (studid, studname, bookname, bookid, dateofissuse, counter_value) from 'C:\Users\Admin\OneDrive\Desktop\BDA Lab\data.csv';
Using 7 child processes

Starting copy of library.library_info with columns [studid, studname, bookname, bookid, dateofissuse, counter_value].
Process ImportProcess-10:      2 rows/s; Avg. rate:      2 rows/s
Traceback (most recent call last):
Process ImportProcess-8:
P Process ImportProcess-11:
Traceback (most recent call last):
Process ImportProcess-9:
P File "C:\Python27\lib\multiprocessing\process.py", line 267, in _bootstrap
  File "C:\Python27\lib\multiprocessing\process.py", line 267, in _bootstrap
Traceback (most recent call last):
Process ImportProcess-10:
P File "C:\Python27\lib\multiprocessing\process.py", line 267, in _bootstrap
  self.run()
Process ImportProcess-12:
Process ImportProcess-14:
Process ImportProcess-13:
  File "C:\apache-cassandra-3.11.13\bin\..\pylib\cqlshlib\copyutil.py", line 2339, in run
    self.run()
Traceback (most recent call last):
  File "C:\Python27\lib\multiprocessing\process.py", line 267, in _bootstrap
Traceback (most recent call last):
  self.run()
Traceback (most recent call last):
  self.run()
File "C:\Python27\lib\multiprocessing\process.py", line 267, in _bootstrap
File "C:\apache-cassandra-3.11.13\bin\..\pylib\cqlshlib\copyutil.py", line 2339, in run
File "C:\Python27\lib\multiprocessing\process.py", line 267, in _bootstrap
  self.close()
  self.run()
File "C:\apache-cassandra-3.11.13\bin\..\pylib\cqlshlib\copyutil.py", line 2339, in run
  self.run()
File "C:\apache-cassandra-3.11.13\bin\..\pylib\cqlshlib\copyutil.py", line 2343, in close
File "C:\apache-cassandra-3.11.13\bin\..\pylib\cqlshlib\copyutil.py", line 2339, in run
  self.close()
File "C:\Python27\lib\multiprocessing\process.py", line 267, in _bootstrap
File "C:\apache-cassandra-3.11.13\bin\..\pylib\cqlshlib\copyutil.py", line 2339, in run
  self.session.cluster.shutdown()
  self.run()
  self.close()
File "C:\apache-cassandra-3.11.13\bin\..\lib\cassandra-driver\internal-only\3.11.0-bb96859b.zip\cassandra-driver\3.11.0-bb96859b\cassandra\cluster.py", line 1259, in shutdown
File "C:\apache-cassandra-3.11.13\bin\..\pylib\cqlshlib\copyutil.py", line 2343, in close
File "C:\apache-cassandra-3.11.13\bin\..\pylib\cqlshlib\copyutil.py", line 2339, in run
  self.close()
  self.close()
File "C:\apache-cassandra-3.11.13\bin\..\pylib\cqlshlib\copyutil.py", line 2343, in close
```

```

File "c:\apache-cassandra-3.11.13\bin\..\\lib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\io\asyncoreactor.py", line 335, in create_timer
File "c:\apache-cassandra-3.11.13\bin\..\\lib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\io\asyncoreactor.py", line 335, in create_timer
File "c:\apache-cassandra-3.11.13\bin\..\\lib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\io\asyncoreactor.py", line 373, in close
self.connection.close()
File "c:\apache-cassandra-3.11.13\bin\..\\lib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\io\asyncoreactor.py", line 335, in create_timer
self.connection.close()
AsyncoreConnection.create_timer(0, partial(asyncore.dispatcher.close, self))
File "c:\apache-cassandra-3.11.13\bin\..\\lib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\io\asyncoreactor.py", line 373, in close
cls.loop.add_timer(timer)
File "c:\apache-cassandra-3.11.13\bin\..\\lib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\io\asyncoreactor.py", line 373, in close
cls.loop.add_timer(timer)
AsyncoreConnection.create_timer(0, partial(asyncore.dispatcher.close, self))
AA AsyncoreConnection.create_timer(0, partial(asyncore.dispatcher.close, self))
AttributeError: 'NoneType' object has no attribute 'add_timer'
File "c:\apache-cassandra-3.11.13\bin\..\\lib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\io\asyncoreactor.py", line 335, in create_timer
AttributeError: 'NoneType' object has no attribute 'add_timer'
File "c:\apache-cassandra-3.11.13\bin\..\\lib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\io\asyncoreactor.py", line 335, in create_timer
AsyncoreConnection.create_timer(0, partial(asyncore.dispatcher.close, self))
cls.loop.add_timer(timer)
cls.loop.add_timer(timer)
A File "c:\apache-cassandra-3.11.13\bin\..\\lib\cassandra-driver-internal-only-3.11.0-bb96859b.zip\cassandra-driver-3.11.0-bb96859b\cassandra\io\asyncoreactor.py", line 335, in create_timer
AttributeError: 'NoneType' object has no attribute 'add_timer'
A cls.loop.add_timer(timer)
cls.loop.add_timer(timer)
AttributeError: 'NoneType' object has no attribute 'add_timer'
AttributeError: 'NoneType' object has no attribute 'add_timer'
AttributeError: 'NoneType' object has no attribute 'add_timer'
Processed: 4 rows; Rate: 1 rows/s; Avg. rate: 2 rows/s
4 rows imported from 1 files in 2.356 seconds (0 skipped).
cqlsh:library>

```

## **MongoDB Lab Program 1 (CRUD Demonstration): -**

### **I. Perform the following DB operations using MongoDB.**

1. Create a database “Student” with the following attributes Rollno, Age, ContactNo, Email-Id.
2. Insert appropriate values
3. Write query to update Email-Id of a student with rollno 10.
4. . Replace the student name from “ABC” to “FEM” of rollno 11

```

Student> db.s.insertMany([
  {StudName:"JohnDoe", Grade:"VIII", Hobbies:"Painting"},
  ... {StudName:"SarahSmith", Grade:"VI", Hobbies:"Reading"},
  ... {StudName:"DavidWilliams", Grade:"IX", Hobbies:"Coding"},
  ... {StudName:"EmilyJones", Grade:"V", Hobbies:"Dancing"},
  ... {StudName:"OliverBrown", Grade:"X", Hobbies:"Photography"},
  ... {StudName:"AvaGarcia", Grade:"IV", Hobbies:"Writing"},
  ... {StudName:"WilliamParker", Grade:"VII", Hobbies:"Chess"},
  ... {StudName:"OliviaWilson", Grade:"III", Hobbies:"Cooking"},
  ... {StudName:"MichaelHall", Grade:"XI", Hobbies:"Gardening"},
  ... {StudName:"CharlotteRoberts", Grade:"II", Hobbies:"Swimming"}])
{
  acknowledged: true,
  insertedIds: {
    '0': ObjectId('660a7cce2d86dcfe9dc00e80'),
    '1': ObjectId('660a7cce2d86dcfe9dc00e81'),
    '2': ObjectId('660a7cce2d86dcfe9dc00e82'),
    '3': ObjectId('660a7cce2d86dcfe9dc00e83'),
    '4': ObjectId('660a7cce2d86dcfe9dc00e84'),
    '5': ObjectId('660a7cce2d86dcfe9dc00e85'),
    '6': ObjectId('660a7cce2d86dcfe9dc00e86'),
    '7': ObjectId('660a7cce2d86dcfe9dc00e87'),
    '8': ObjectId('660a7cce2d86dcfe9dc00e88'),
    '9': ObjectId('660a7cce2d86dcfe9dc00e89')
  }
}
Student> db.s.find()
[
  {
    _id: ObjectId('660a7a7e2d86dcfe9dc00e7e'),
    StudName: 'MichelleJacintha',
    Grade: 'VII',
    Hobbies: 'InternsetSurfing'
  },
  {
    _id: ObjectId('660a7ab22d86dcfe9dc00e7f'),
    StudName: 'Aryana',
    Grade: 'VIII',
    Hobbies: 'InternsetSurfing'
  },
  {
    _id: ObjectId('660a7cce2d86dcfe9dc00e80'),
    StudName: 'JohnDoe',
    Grade: 'VIII',
    Hobbies: 'Painting'
  },
  {
    _id: ObjectId('660a7cce2d86dcfe9dc00e81'),
    StudName: 'SarahSmith',
    Grade: 'VI',
    Hobbies: 'Reading'
  },
  {
    _id: ObjectId('660a7cce2d86dcfe9dc00e82'),
    StudName: 'DavidWilliams',
    Grade: 'IX',
    Hobbies: 'Coding'
  },
  {
    _id: ObjectId('660a7cce2d86dcfe9dc00e83'),
    StudName: 'EmilyJones',
    Grade: 'V',
    Hobbies: 'Dancing'
  },
  {
    _id: ObjectId('660a7cce2d86dcfe9dc00e84'),
    StudName: 'OliverBrown',
    Grade: 'X',
    Hobbies: 'Photography'
  },
  {
    _id: ObjectId('660a7cce2d86dcfe9dc00e85'),
    StudName: 'AvaGarcia',
    Grade: 'IV',
    Hobbies: 'Writing'
  },
  {
    _id: ObjectId('660a7cce2d86dcfe9dc00e86'),
    StudName: 'WilliamParker',
    Grade: 'VII',
    Hobbies: 'Chess'
  },
  {
    _id: ObjectId('660a7cce2d86dcfe9dc00e87'),
    StudName: 'OliviaWilson',
    Grade: 'III',
    Hobbies: 'Cooking'
  },
  {
    _id: ObjectId('660a7cce2d86dcfe9dc00e88'),
    StudName: 'MichaelHall',
    Grade: 'XI',
    Hobbies: 'Gardening'
  },
  {
    _id: ObjectId('660a7cce2d86dcfe9dc00e89'),
    StudName: 'CharlotteRoberts',
    Grade: 'II',
    Hobbies: 'Swimming'
  }
]

```



```

    Grade: 'III',
    Hobbies: 'Cooking'
  },
  {
    _id: ObjectId('660a7cce2d86dcfefdc00e88'),
    StudName: 'MichaelHall',
    Grade: 'XI',
    Hobbies: 'Gardening'
  },
  {
    _id: ObjectId('660a7cce2d86dcfefdc00e89'),
    StudName: 'CharlotteRoberts',
    Grade: 'II',
    Hobbies: 'Swimming'
  }
]
Student> db.s.upodateOne({StudName:"AvaGarcia"}, {$set:{StudName: 'Ambika'}})
TypeError: db.s.upodateOne is not a function
Student> db.s.updateOne({StudName:"AvaGarcia"}, {$set:{StudName: 'Ambika'}})
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
Student> db.s.find({StudName:'Ambika'})
[
  {
    _id: ObjectId('660a7cce2d86dcfefdc00e85'),
    StudName: 'Ambika',
    Grade: 'IV',
    Hobbies: 'Writing'
  }
]
Student> db.s.updateOne({StudName:"Ada"}, {$set:{StudName: 'Ambuja'}}, {upsert: true})
{
  acknowledged: true,
  insertedId: ObjectId('660a7e7c7b7d50080592d516'),
  matchedCount: 0,
  modifiedCount: 0,
  upsertedCount: 1
}
Student> db.s.find({StudName:'Ambuja'})
[ { _id: ObjectId('660a7e7c7b7d50080592d516'), StudName: 'Ambuja' } ]
Student> db.s.updateOne({StudName:"Ambuja"}, {$set:{Grade: 'LKG'}, {Hobbies: 'Staring at Wall'}}, {upsert: true})
Uncaught:
SyntaxError: Unexpected token (1:58)

> 1 | db.s.updateOne({StudName:"Ambuja"}, {$set:{Grade: 'LKG'}, {Hobbies: 'Staring at Wall'}}, {upsert: true})
    |                                     ^
    2 |

Student> db.s.updateOne({StudName:"Ambuja"}, {$set:{Grade: 'LKG', Hobbies: 'Staring at Wall'}}, {upsert: true})
{
  acknowledged: true,
  insertedId: ObjectId('660a7e7c7b7d50080592d516'),
  matchedCount: 0,
  modifiedCount: 0,
  upsertedCount: 1
}

```

```

}
]
Student> db.s.updateOne({StudName:"Ada"}, {$set:{StudName: 'Ambuja'}}, {upsert: true})
{
  acknowledged: true,
  insertedId: ObjectId('660a7e7c7b7d50080592d516'),
  matchedCount: 0,
  modifiedCount: 0,
  upsertedCount: 1
}
Student> db.s.find({StudName:'Ambuja'})
[ { _id: ObjectId('660a7e7c7b7d50080592d516'), StudName: 'Ambuja' } ]
Student> db.s.updateOne({StudName:"Ambuja"}, {$set:{Grade: 'LKG'}, {Hobbies: 'Staring at Wall'}}, {upsert: true})
Uncaught:
SyntaxError: Unexpected token (1:58)

> 1 | db.s.updateOne({StudName:"Ambuja"}, {$set:{Grade: 'LKG'}, {Hobbies: 'Staring at Wall'}}, {upsert: true})
    |                                     ^
    2 |

Student> db.s.updateOne({StudName:"Ambuja"}, {$set:{Grade: 'LKG', Hobbies: 'Staring at Wall'}}, {upsert: true})
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
Student> db.s.find({StudName:'Ambuja'})
[
  {
    _id: ObjectId('660a7e7c7b7d50080592d516'),
    StudName: 'Ambuja',
    Grade: 'LKG',
    Hobbies: 'Staring at Wall'
  }
]
Student> db.s.updateOne({StudName:"Ambuja"}, {$set:{Grade: 'LKG', Hobbies: 'Chewing Gum'}}, {upsert: true})
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
Student> db.s.find({StudName:'Ambuja'})
[
  {
    _id: ObjectId('660a7e7c7b7d50080592d516'),
    StudName: 'Ambuja',
    Grade: 'LKG',
    Hobbies: 'Chewing Gum'
  }
]
Student> 

```

## II. Perform the following DB operations using MongoDB.

1. Create a collection by name Customers with the following attributes.  
Cust\_id, Acc\_Bal, Acc\_Type
2. Insert at least 5 values into the table
3. Write a query to display those records whose total account balance is greater than 1200 of account type 'Z' for each customer\_id.
4. Determine Minimum and Maximum account balance for each customer\_i

```

Student> db.createCollection("c")
{ ok: 1 }
Student> db.c.insertMany([
...   {Cust_id: "0001", Acc_Bal: 5000, Acc_type: "Savings"},
...   {Cust_id: "0002", Acc_Bal: 12000, Acc_type: "Current"},
...   {Cust_id: "0003", Acc_Bal: 7500, Acc_type: "Savings"},
...   {Cust_id: "0004", Acc_Bal: 25000, Acc_type: "Current"},
...   {Cust_id: "0005", Acc_Bal: 10000, Acc_type: "Savings"},
...   {Cust_id: "0006", Acc_Bal: 17000, Acc_type: "Current"},
...   {Cust_id: "0007", Acc_Bal: 8000, Acc_type: "Savings"},
...   {Cust_id: "0008", Acc_Bal: 32000, Acc_type: "Current"},
...   {Cust_id: "0009", Acc_Bal: 15000, Acc_type: "Savings"},
...   {Cust_id: "0010", Acc_Bal: 20000, Acc_type: "Current"}
... ])
{
  acknowledged: true,
  insertedIds: {
    '0': ObjectId('660a84152d86dcfe9dc00e8a'),
    '1': ObjectId('660a84152d86dcfe9dc00e8b'),
    '2': ObjectId('660a84152d86dcfe9dc00e8c'),
    '3': ObjectId('660a84152d86dcfe9dc00e8d'),
    '4': ObjectId('660a84152d86dcfe9dc00e8e'),
    '5': ObjectId('660a84152d86dcfe9dc00e8f'),
    '6': ObjectId('660a84152d86dcfe9dc00e90'),
    '7': ObjectId('660a84152d86dcfe9dc00e91'),
    '8': ObjectId('660a84152d86dcfe9dc00e92'),
    '9': ObjectId('660a84152d86dcfe9dc00e93')
  }
}
Student> 

```

```

Student> db.c.find({Acc_Bal: {$gt: 1200}, Acc_type: "Current"})
[
  {
    _id: ObjectId('660a84152d86dcfe9dc00e8b'),
    Cust_id: '0002',
    Acc_Bal: 12000,
    Acc_type: 'Current'
  },
  {
    _id: ObjectId('660a84152d86dcfe9dc00e8d'),
    Cust_id: '0004',
    Acc_Bal: 25000,
    Acc_type: 'Current'
  },
  {
    _id: ObjectId('660a84152d86dcfe9dc00e8f'),
    Cust_id: '0006',
    Acc_Bal: 17000,
    Acc_type: 'Current'
  },
  {
    _id: ObjectId('660a84152d86dcfe9dc00e91'),
    Cust_id: '0008',
    Acc_Bal: 32000,
    Acc_type: 'Current'
  },
  {
    _id: ObjectId('660a84152d86dcfe9dc00e93'),
    Cust_id: '0010',
    Acc_Bal: 20000,
    Acc_type: 'Current'
  }
]
Student> 

```

```

Student> db.c.aggregate([{"$group": {"_id": null, "max": {"$max": "$Acc_Bal"}, "min": {"$min": "$Acc_Bal" }}}])
[ { _id: null, max: 32000, min: 5000 } ]
Student> db.c.aggregate([{"$group": {"_id": "Cust_id", "max": {"$max": "$Acc_Bal"}, "min": {"$min": "$Acc_Bal" }}}])
[ { _id: 'Cust_id', max: 32000, min: 5000 } ]
Student> 

```

## Hadoop Commands

```
hdusersbmsce-OptiPlus-3000:~$ sudo su hduser
```

```
[sudo] password for hduser:
```

```
hdusersbmsce-OptiPlus-3000:~$ start-all.sh
```

This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh

```
22/06/06 14:43:45 WARN util.NativeCodeLoader: Unable to load
native-hadoop Library for your platform... using builtin-java classes where
applicable Starting namenodes on [localhost] localhost: namenode running as
process 3396. Stop it first. localhost: datanode running as process 3564, Stop it
first.
```

```
starting secondary namenodes [0.0.0.0] 0.0.0.0: secondarynamenode running as process 3773. Stop it first.
022/06/06 14:43:47 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your starting
yarn daemons resource process 3932. Stop it first.
```

```
Localhost: running as process 4255. stop it
first. 6003 Jps
```

```
3932 ResourceManager
```

```
3773 SecondaryNameNode
```

```
4255 NodeManager
```

```
hdusersbmsce-OptiPlus-3060:~$ hdfs dfs -mkdir
```

```
/Amrutha hdusersbmsce-OptiPlus-3060:~$ hdfs dfs -ls /
```

```
22/06/06 14:45:30 WARN util.NativeCodeLoader: Unable to load
native-hadoop library for your platform... using builtin-java classes where
applicable Found 19 items
```

```
drwxr-xr-x          hduser
```

```
supergroup 02022- 06-06
```

```
11:44 /AAA
```

```
drwxr-xr-x          -hduser
```

```
supergroup 2022-06-03 12:17
```

```
/Army drwxr-xr-x      hduser
```

```
supergroup 02022- 06-06 11:40
```

```
/Avnit drwxr-xr-x      -hduser
```

```
supergroup 02022-05-31 10:44 /88
```

```
drwxr-xr-x -hduser supergroup
```

```
02022- 06-01 15:03 /Cath
```

```
drwxr-xr-x -hduser supergroup
```

```
drwxr-xr-x hduser supergroup
```

```
drwxr-xr-x -hduser supergroup
```

```
drwxr-xr-x -hduser supergroup
```

```
drwxr-xr-x -hduser supergroup
```

```
drwxr-xr-x -hduser supergroup
```

```
drwxr-xr-x -hduser supergroup
```

```
drwxr-xr-x -hduser supergroup
```

```
drwxr-xr-x -hduser supergroup
```

```
drwxr-xr-x - hduser
```



```
supergroup drwxr-xr-x -hduser
supergroup 82022-06-04 10:06
/FFF
02022-06-06 14:40 /Kmr
02022-06-06 14:44 /Amrutha
02022-06-01 15:03 /Neha
02022-06-04 09:54 /WC.txt
0 2022-06-04 09:54 /welcome.txt
02022-06-06 11:36 /abc
62022-06-03 12:13 /akash
0 2022-06-03 15:12 /darshan
0 2022-06-04 09:31 /ghh 8 2022-
06-06 11:45 /hello drwxr-xr-x -
hduser supergroup
62022-06-04 09:35 /rahul
drwxr-xr-x -hduser
supergroup 02022-06-03 12:11
/shre drwxr-xr-x .hduser supergroup 02022-06-03
12:41
/shreshtha
hdusersbmsce-OptiPlus-3060:-$ hdfs dfs put /home/hduser/Desktop/6b.txt
/Amrutha/WC.txt
22/05/06 14:46:40 WARN util.NativeCodeLoader: Unable to load
native-hadoop library for your platform... using builtin-java classes where
applicable hduserabesce-OptiPlex-3060:-$ hdfs dfs cat /Amrutha/WC.txt
22/06/06 14:47:00 WARN util.NativeCodeLoader: Unable to load native-hadoop
library for your platform... using builtin-java classes where applicable hello from
of hdusersbmsce-OptiPlus-3040:-$ hdfs dfs-get /Amrutha/WC.txt
/home/hduser/Downloads/newic.txt
22/05/06 14:51:43 WARN util.NativeCodeLoader: Unable to load native-hadoop
library for your platform... using builtin-java classes where applicable
hdusersbmsce-OptiPlus-3066:-$ cd Downloads
hdusersbmsce-OptiPlus-3060:-/Downloads$ cat newwMC.Ext hello from 6E
hdusersbmsce-OptiPlus-3060:-$ hdfs dfs -ls /Amrutha/
22/06/06 14:54:04 WARN util.NativeCodeLoader: Unable to load
native-hadoop Library for your platform... using builtin java classes where
applicable
Found 2 items
-rw-r--r-- 1 hduser supergroup
23 2022-06-06 14:46 /Amrutha/MC.txt
1 hduser supergroup
23 2022-06-06 14:58 /Amrutha/newwc.txt
hdusersbmsce-OptiPlus-3060:-$ hdfs dfs -getmerge /Amrutha/wc.txt
/Amrutha/newwc.txt /home/hduser/Desktop/newmerge.txt
```

22/06/06 14:55:18 NARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-Java classes where applicable  
hduser@besce-OptiPlex-3060:~\$ cd Desktop

hduser@besce-OptiPlex-3060:~/Desktops cat newmerge.txt hello from 68

D B

hello

from 68

D B

hdusersbmsce-OptiPlus-3060:~/Desktops hadoop fs getfacl /Amrutha/ 22/06/06

14:56:24 WARN util.NativeCodeLoader: Unable to load native hadoop library for your platform... using builtin java classes where applicable # file: /Amrutha

# owner: hduser #

group:

supergroup

user::rwx

group::r-x

other::r-x

hdusersbmsce-OptiPlus-3060:~/Desktop5 hdfs dfs copyToLocal /Amrutha/HC.txt

/home/hduser/Desktop

22/05/06 14:58:09 WARN util.NativeCodeLoader: Unable to load native-hadoop Library for your platform... using builtin-java classes where applicable

hdusersbmsce-OptiPlus-3000:-

/Desktop5 cat MC.txt hello from 68

hdusersbmsce-OptiPlus-3060:~/Desktops hdfs dfs -cat /Amrutha/MC.txt 22/06/06 14:58:59 WARN util.NativeCodeLoader: Unable to load native-hadoop Library for your platform... using builtin-Java classes where applicable hello from GB B

hdusersbmsce-OptiPlus-3060:~/Desktop5 hadoop fs - /Amrutha /FFF 22/06/06

14:59:46 WARN util.NativeCodeLoader: Unable to load native-hadoop Library for your platform... using builtin-java classes where applicable

hduser@besce-OptiPlex-3060:~/Desktops hadoop fs -Ls /FFF 22/05/06 15:00:00 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform...

using builtin- java classes where applicable Found 2 items drwxr-xr-x -hduser supergroup TWEE 1 hduser supergroup 02022-05-06 14:50

/FFF/Amrutha 17 2022-05-04 10:06 /FFF/MC.txt

hdusersbmsce-OptiPlus-3060:~/Desktops hadoop fs cp /FFF/ /LLL

22/06/06 15:09:34 WARN util.NativeCodeLoader: Unable to load native hadoop library for your platform... using builtin-java classes where applicable

hdusersbmsce-OptiPlus-3060:-

/Desktops hadoop fs -Ls /LLL

22/06/06 15:10:07 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable Found 2 items

drwxr-xr-x -hduser supergroup

hdusersbmsce-OptiPlus-3000:~/Desktops 02022-06-06 15:09

## Hadoop Programs

1) Word Count WCMapper Java Class file.

```
// Importing libraries import java.io.IOException; import org.apache.hadoop.io.IntWritable; import
org.apache.hadoop.io.LongWritable; import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase; import
org.apache.hadoop.mapred.Mapper; import
org.apache.hadoop.mapred.OutputCollector; import
org.apache.hadoop.mapred.Reporter;
```

```
public class WCMapper extends MapReduceBase implements Mapper<LongWritable,
    Text, Text, IntWritable> {
```

```
    // Map function
```

```
    public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output,
Reporter rep) throws IOException
    {
```

```
        String line = value.toString();
```

```
        // Splitting the line on spaces
for (String word : line.split(" "))
```

```
    {
        if (word.length() > 0)
        {
            output.collect(new Text(word), new IntWritable(1));
        } } }
```

Reducer Code

```
// Importing libraries import java.io.IOException; import
java.util.Iterator; import
org.apache.hadoop.io.IntWritable; import
```

```

org.apache.hadoop.io.Text;  import  org.apache.hadoop.mapred.MapReduceBase;  import
org.apache.hadoop.mapred.OutputCollector;  import  org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;

public class WCReducer extends MapReduceBase implements Reducer<Text, IntWritable, Text,
                                IntWritable> {

    // Reduce function  public  void  reduce(Text  key,
    Iterator<IntWritable> value,
                                OutputCollector<Text,  IntWritable>  output,
                                Reporter rep) throws IOException
    {

        int count = 0;

        // Counting the frequency of each words
        while (value.hasNext())
        {
            IntWritable i = value.next();
            count += i.get();
        }

        output.collect(key, new IntWritable(count));
    }
}

```

```

Driver Code: // Importing libraries import java.io.IOException;
import      org.apache.hadoop.conf.Configured;          import
org.apache.hadoop.fs.Path;                               import
org.apache.hadoop.io.IntWritable; import

```



```

org.apache.hadoop.io.Text;
org.apache.hadoop.mapred.FileInputFormat;
org.apache.hadoop.mapred.FileOutputFormat;
org.apache.hadoop.mapred.JobClient;
org.apache.hadoop.mapred.JobConf;
org.apache.hadoop.util.Tool;
org.apache.hadoop.util.ToolRunner;
public class WCDriver extends Configured implements Tool {

```

```

import
import
import
import
import
import

```

```

    public int run(String args[]) throws IOException

```

```

    {
        if (args.length < 2)
        {
            System.out.println("Please give valid inputs"); return -1;
        }

```

```

        JobConf conf = new JobConf(WCDriver.class); FileInputFormat.setInputPaths(conf, new
        Path(args[0]));
        FileOutputFormat.setOutputPath(conf, new Path(args[1]));
        conf.setMapperClass(WCMapper.class);
        conf.setReducerClass(WCReducer.class);
        conf.setMapOutputKeyClass(Text.class);
        conf.setMapOutputValueClass(IntWritable.class);
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        JobClient.runJob(conf)
        ; return 0;
    }

```

```

// Main Method

```

```

public static void main(String args[]) throws Exception{

```

```

    int    exitCode    =    ToolRunner.run(new    WCDriver(),    args);
    System.out.println(exitCode);
}
} Output :

```

```

hduser@bmsce-Precision-T1700:~$ su hduser\
> ^C
hduser@bmsce-Precision-T1700:~$ ^C
hduser@bmsce-Precision-T1700:~$ su hduser
Password:
hduser@bmsce-Precision-T1700:~$ start-all.sh
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
Starting namenodes on [localhost]
hduser@localhost's password:
localhost: starting namenode, logging to /usr/local/hadoop/logs/hadoop-hduser-namenode-bmsce-
Precision-T1700.out
hduser@localhost's password:
localhost: starting datanode, logging to /usr/local/hadoop/logs/hadoop-hduser-datanode-bmsce-
Precision-T1700.out
Starting secondary namenodes [0.0.0.0]
hduser@0.0.0.0's password:
0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop/logs/hadoop-hduser-
secondarynamenode-bmsce-Precision-T1700.out
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop/logs/yarn-hduser-resourcemanager-bmsce-
Precision-T1700.out
hduser@localhost's password:
localhost: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-hduser-nodemanager-
bmsce-Precision-T1700.out
hduser@bmsce-Precision-T1700:~$ jps
7328 Jps
6497 DataNode
4372 org.eclipse.equinox.launcher_1.5.600.v20191014-2022.jar
6325 NameNode
7206 NodeManager
6872 ResourceManager
6713 SecondaryNameNode
hduser@bmsce-Precision-T1700:~$ cat > sample.txt
hi im khushil
i am learing hadoop
hadoop is awesome
^C
hduser@bmsce-Precision-T1700:~$ cat sample.txt
hi im khushil
i am learing hadoop
hadoop is awesome
hduser@bmsce-Precision-T1700:~$ hdfs dfs -ls /
Found 18 items
drwxr-xr-x - hduser supergroup 0 2022-06-06 12:35 /CSE
drwxr-xr-x - hduser supergroup 0 2022-06-06 12:23 /FFF
drwxr-xr-x - hduser supergroup 0 2022-06-06 12:36 /LLL
drwxr-xr-x - hduser supergroup 0 2022-06-20 12:06 /amit_bda
drwxr-xr-x - hduser supergroup 0 2022-06-03 14:52 /bharath
drwxr-xr-x - hduser supergroup 0 2022-06-03 14:43 /bharath035
drwxr-xr-x - hduser supergroup 0 2022-05-31 10:21 /example
drwxr-xr-x - hduser supergroup 0 2022-06-01 15:13 /foldernew
drwxr-xr-x - hduser supergroup 0 2022-06-06 15:04 /hemang061
drwxr-xr-x - hduser supergroup 0 2022-06-03 12:27 /irfan
drwxr-xr-x - hduser supergroup 0 2022-06-01 15:09 /muskan
drwxr-xr-x - hduser supergroup 0 2022-06-06 15:04 /new_folder
drwxr-xr-x - hduser supergroup 0 2022-05-31 10:26 /one
drwxr-xr-x - hduser supergroup 0 2022-06-20 12:17 /output
drwxr-xr-x - hduser supergroup 0 2022-06-03 12:08 /saurab
drwxrwxr-x - hduser supergroup 0 2019-08-01 16:19 /tmp
drwxr-xr-x - hduser supergroup 0 2019-08-01 16:03 /user

```

```

drwxr-xr-x - hduser supergroup 0 2022-06-01 09:46 /user1
hduser@bmsce-Precision-T1700:~$ hdfs dfs -mkdir /input_khushil
hduser@bmsce-Precision-T1700:~$ hdfs dfs -ls /
Found 19 items
drwxr-xr-x - hduser supergroup 0 2022-06-06 12:35 /CSE
drwxr-xr-x - hduser supergroup 0 2022-06-06 12:23 /FFF
drwxr-xr-x - hduser supergroup 0 2022-06-06 12:36 /LLL
drwxr-xr-x - hduser supergroup 0 2022-06-20 12:06 /amit_bda
drwxr-xr-x - hduser supergroup 0 2022-06-03 14:52 /bharath
drwxr-xr-x - hduser supergroup 0 2022-06-03 14:43 /bharath035
drwxr-xr-x - hduser supergroup 0 2022-05-31 10:21 /example
drwxr-xr-x - hduser supergroup 0 2022-06-01 15:13 /foldernew
drwxr-xr-x - hduser supergroup 0 2022-06-06 15:04 /hemang061
drwxr-xr-x - hduser supergroup 0 2022-06-20 15:13 /input_khushil
drwxr-xr-x - hduser supergroup 0 2022-06-03 12:27 /irfan
drwxr-xr-x - hduser supergroup 0 2022-06-01 15:09 /muskan
drwxr-xr-x - hduser supergroup 0 2022-06-06 15:04 /new_folder
drwxr-xr-x - hduser supergroup 0 2022-05-31 10:26 /one
drwxr-xr-x - hduser supergroup 0 2022-06-20 12:17 /output
drwxr-xr-x - hduser supergroup 0 2022-06-03 12:08 /saurab
drwxrwxr-x - hduser supergroup 0 2019-08-01 16:19 /tmp
drwxr-xr-x - hduser supergroup 0 2019-08-01 16:03 /user
drwxr-xr-x - hduser supergroup 0 2022-06-01 09:46 /user1
hduser@bmsce-Precision-T1700:~$ hdfs dfs -put /home/hduser/sample.txt /input_khushil
hduser@bmsce-Precision-T1700:~$ hdfs dfs -ls /input_khushil
Found 1 items
-rw-r--r-- 1 hduser supergroup 52 2022-06-20 15:15 /input_khushil/sample.txt
hduser@bmsce-Precision-T1700:~$ hadoop jar /home/hduser/khushil/WordCount.jar WCDriver
/input_khushil /input_khushil/output_khushil
22/06/20 15:16:44 INFO Configuration.deprecation: session.id is deprecated. Instead, use
dfs.metrics.session-id
22/06/20 15:16:44 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker,
sessionId=
22/06/20 15:16:44 INFO jvm.JvmMetrics: Cannot initialize JVM Metrics with
processName=JobTracker, sessionId= - already initialized
22/06/20 15:16:44 WARN mapreduce.JobSubmitter: Hadoop command-line option parsing not
performed. Implement the Tool interface and execute your application with ToolRunner to remedy
this.
22/06/20 15:16:44 INFO mapred.FileInputFormat: Total input paths to process : 1
22/06/20 15:16:44 INFO mapreduce.JobSubmitter: number of splits:1
22/06/20 15:16:44 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_local230197290_0001
22/06/20 15:16:44 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
22/06/20 15:16:44 INFO mapred.LocalJobRunner: OutputCommitter set in config null
22/06/20 15:16:44 INFO mapreduce.Job: Running job: job_local230197290_0001
22/06/20 15:16:44 INFO mapred.LocalJobRunner: OutputCommitter is
org.apache.hadoop.mapred.FileOutputCommitter
22/06/20 15:16:44 INFO mapred.LocalJobRunner: Waiting for map tasks
22/06/20 15:16:44 INFO mapred.LocalJobRunner: Starting task:
attempt_local230197290_0001_m_000000_0
22/06/20 15:16:44 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
22/06/20 15:16:44 INFO mapred.MapTask: Processing split:
hdfs://localhost:54310/input_khushil/sample.txt:0+52
22/06/20 15:16:44 INFO mapred.MapTask: numReduceTasks: 1
22/06/20 15:16:44 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
22/06/20 15:16:44 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
22/06/20 15:16:44 INFO mapred.MapTask: soft limit at 83886080
22/06/20 15:16:44 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
22/06/20 15:16:44 INFO mapred.MapTask: kvstart = 26214396; length = 6553600

```

```

GC time elapsed (ms)=1
CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=471859200

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=52
File Output Format Counters
  Bytes Written=63
0
hduser@bmsce-Precision-T1700:~$ hdfs dfs -ls /input_khushil
Found 2 items
drwxr-xr-x  - hduser supergroup          0 2022-06-20 15:16 /input_khushil/output_khushil
-rw-r--r--  1 hduser supergroup          52 2022-06-20 15:15 /input_khushil/sample.txt
hduser@bmsce-Precision-T1700:~$ hdfs dfs -ls /input_khushil/output_khushil
Found 2 items
-rw-r--r--  1 hduser supergroup          0 2022-06-20 15:16
/input_khushil/output_khushil/_SUCCESS

```

```

-rw-r--r--  1 hduser supergroup          63 2022-06-20 15:16
/input_khushil/output_khushil/part-00000
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cat /input_khushil/output_khushil/part-0000
cat: '/input_khushil/output_khushil/part-0000': No such file or directory
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cat /input_khushil/output_khushil/part-00000
am      1
awesome 1
hadoop 2
hi      1
i       1
im      1
is      1
khushil 1
learing 1

```

2) Top N  
Driver-TopN.class

**package** samples.topn;



```

import java.io.IOException;      import java.util.StringTokenizer;      import
org.apache.hadoop.conf.Configuration;      import org.apache.hadoop.fs.Path;      import
org.apache.hadoop.io.IntWritable;      import org.apache.hadoop.io.Text;      import
org.apache.hadoop.mapreduce.Job; import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import
org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.GenericOptionsParser;

```

```

public class TopN {                public static void main(String[] args) throws
Exception {
    Configuration conf = new Configuration();
    String[] otherArgs = (new GenericOptionsParser(conf, args)).getRemainingArgs();
                                if (otherArgs.length != 2) {
        System.err.println("Usage: TopN <in> <out>"); System.exit(2);
    }
    Job job = Job.getInstance(conf); job.setJobName("Top N");
job.setJarByClass(TopN.class);
job.setMapperClass(TopNMapper.class);
job.setReducerClass(TopNReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(otherArgs[0])); FileOutputFormat.setOutputPath(job, new
Path(otherArgs[1])); System.exit(job.waitForCompletion(true) ? 0 : 1);
}    public static class TopNMapper extends Mapper<Object, Text, Text, IntWritable> {
    private static final IntWritable one = new IntWritable(1);

    private Text word = new Text();

```

```

        private String tokens =
        "[_!$#<>\\^=\\[\\]\\*\\/\\\\,;,.\\:()?!\\\"'"]";
        public void map(Object key, Text value, Mapper<Object, Text, Text,
        IntWritable>.Context context) throws IOException,
        InterruptedException {
            String cleanLine =
            value.toString().toLowerCase().replaceAll(this.tokens, " ");
            StringTokenizer itr = new
            StringTokenizer(cleanLine);
            while (itr.hasMoreTokens()) {
                this.word.set(itr.nextToken().trim());
                context.write(this.word, one);
            }
        }
    }
}

```

TopNCombiner.class  
**package** samples.topn;

```

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

```

```

public class TopNCombiner extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Reducer<Text, IntWritable, Text,
    IntWritable>.Context context) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) sum += val.get();
        context.write(key, new IntWritable(sum));
    }
}

```

TopNMapper.class  
**package** samples.topn;

```

import java.io.IOException; import java.util.StringTokenizer;
import org.apache.hadoop.io.IntWritable; import
org.apache.hadoop.io.Text; import
org.apache.hadoop.mapreduce.Mapper;

public class TopNMapper extends Mapper<Object, Text, Text, IntWritable> { private static
final IntWritable one = new IntWritable(1);
    private Text word = new
Text();
    private String tokens =
"[_|$.#<\\^=\\[\\]\\*\\/\\\\\\,;,.\\:\\()?!\\\"'"]";
    public void map(Object key, Text value, Mapper<Object, Text, Text,
IntWritable>.Context context) throws IOException,
InterruptedException { String cleanLine =
value.toString().toLowerCase().replaceAll(this.tokens, " "); StringTokenizer itr = new
StringTokenizer(cleanLine);
    while (itr.hasMoreTokens()) {
    this.word.set(itr.nextToken().trim());
context.write(this.word, one);
    }
}
}

```

TopNReducer.class

```
package samples.topn;
```

```

import java.io.IOException; import
java.util.HashMap; import java.util.Map; import
org.apache.hadoop.io.IntWritable; import
org.apache.hadoop.io.Text; import
org.apache.hadoop.mapreduce.Reducer; import
utils.MiscUtils;

```

```

public class TopNReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
private Map<Text, IntWritable> countMap = new HashMap<>();

    public void reduce(Text key, Iterable<IntWritable> values, Reducer<Text, IntWritable, Text,
IntWritable>.Context context) throws IOException, InterruptedException {
        int
sum = 0; for (IntWritable val : values)
            sum += val.get();
this.countMap.put(new Text(key), new IntWritable(sum)); }

    protected void cleanup(Reducer<Text, IntWritable, Text, IntWritable>.Context
context) throws IOException,
InterruptedException {
        Map<Text, IntWritable> sortedMap =
MiscUtils.sortByValues(this.countMap); int counter = 0; for
(Text key : sortedMap.keySet()) {
            if (counter++
== 20)
                break;
            context.write(key, sortedMap.get(key));
        }
    }
}

```

Output:

```

hduser@bmsce-Precision-T1700:~/Desktop/temperature$ hdfs dfs -mkdir /khushil_topn
hduser@bmsce-Precision-T1700:~/Desktop/temperature$ hdfs dfs -put ./input.txt /khushil_topn/
hduser@bmsce-Precision-T1700:~/Desktop/temperature$ hdfs dfs -ls /khushil_topn/
Found 1 items
-rw-r--r--  1 hduser supergroup      103 2022-06-27 15:43 /khushil_topn/input.txt
hduser@bmsce-Precision-T1700:~/Desktop/temperature$ hadoop jar topn.jar TopNDriver
/khushil_topn/input.txt /khushil_topn/output
Exception in thread "main" java.lang.ClassNotFoundException: TopNDriver
    at java.net.URLClassLoader.findClass(URLClassLoader.java:382)
    at java.lang.ClassLoader.loadClass(ClassLoader.java:418)
    at java.lang.ClassLoader.loadClass(ClassLoader.java:351)
    at java.lang.Class.forName0(Native Method)
    at java.lang.Class.forName(Class.java:348)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:214)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
hduser@bmsce-Precision-T1700:~/Desktop/temperature$ hadoop jar topn.jar topn.TopNDriver
/khushil_topn/input.txt /khushil_topn/output
22/06/27 15:45:22 INFO Configuration.deprecation: session.id is deprecated. Instead, use
dfs.metrics.session-id
22/06/27 15:45:22 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker,
sessionId=
22/06/27 15:45:22 INFO input.FileInputFormat: Total input paths to process : 1
22/06/27 15:45:22 INFO mapreduce.JobSubmitter: number of splits:1
22/06/27 15:45:22 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local691635730_0001
22/06/27 15:45:22 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
22/06/27 15:45:22 INFO mapreduce.Job: Running job: job_local691635730_0001
22/06/27 15:45:22 INFO mapred.LocalJobRunner: OutputCommitter set in config null
22/06/27 15:45:22 INFO mapred.LocalJobRunner: OutputCommitter is
org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
22/06/27 15:45:22 INFO mapred.LocalJobRunner: Waiting for map tasks
22/06/27 15:45:22 INFO mapred.LocalJobRunner: Starting task: attempt_local691635730_0001_m_000000_0
22/06/27 15:45:22 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
22/06/27 15:45:22 INFO mapred.MapTask: Processing split:
hdfs://localhost:54310/khushil_topn/input.txt:0+103
22/06/27 15:45:22 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
22/06/27 15:45:22 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
22/06/27 15:45:22 INFO mapred.MapTask: soft limit at 83886080
22/06/27 15:45:22 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
22/06/27 15:45:22 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
22/06/27 15:45:22 INFO mapred.MapTask: Map output collector class =
org.apache.hadoop.mapred.MapTask$MapOutputBuffer
22/06/27 15:45:22 INFO mapred.LocalJobRunner:
22/06/27 15:45:22 INFO mapred.MapTask: Starting flush of map output
22/06/27 15:45:22 INFO mapred.MapTask: Spilling map output
22/06/27 15:45:22 INFO mapred.MapTask: bufstart = 0; bufend = 187; bufvoid = 104857600
22/06/27 15:45:22 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214316(104857264);
length = 81/6553600
22/06/27 15:45:22 INFO mapred.MapTask: Finished spill 0
22/06/27 15:45:22 INFO mapred.Task: Task:attempt_local691635730_0001_m_000000_0 is done. And is in
the process of committing
22/06/27 15:45:22 INFO mapred.LocalJobRunner: map
22/06/27 15:45:22 INFO mapred.Task: Task 'attempt_local691635730_0001_m_000000_0' done.
22/06/27 15:45:22 INFO mapred.LocalJobRunner: Finishing task: attempt_local691635730_0001_m_000000_0
22/06/27 15:45:22 INFO mapred.LocalJobRunner: map task executor complete.
22/06/27 15:45:22 INFO mapred.LocalJobRunner: Waiting for reduce tasks
22/06/27 15:45:22 INFO mapred.LocalJobRunner: Starting task: attempt_local691635730_0001_r_000000_0
22/06/27 15:45:22 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]

```



```

Map input records=6
Map output records=21
Map output bytes=187
Map output materialized bytes=235
Input split bytes=110
Combine input records=0
Combine output records=0
Reduce input groups=15
Reduce shuffle bytes=235
Reduce input records=21
Reduce output records=15
Spilled Records=42
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=42
CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=578289664
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=103
File Output Format Counters
Bytes Written=105
hduser@bmsce-Precision-T1700:~/Desktop/temperature$ hdfs dfs -ls /khushil_topn/output/
Found 2 items
-rw-r--r--  1 hduser supergroup          0 2022-06-27 15:45 /khushil_topn/output/_SUCCESS
-rw-r--r--  1 hduser supergroup       105 2022-06-27 15:45 /khushil_topn/output/part-r-00000
hduser@bmsce-Precision-T1700:~/Desktop/temperature$ hdfs dfs -cat /khushil_topn/output/part-r-00000
hadoop  4
i3
am      2
hi      1
im      1
is      1
there   1
bye     1
learing 1
awesome 1
love    1
khushil 1
cool    1
and     1
using   1
hduser@bmsce-Precision-T1700:~/Desktop/temperature$

```

### 3) Average Temperature



## AverageDriver

**package** temp;

```
import    org.apache.hadoop.fs.Path;    import    org.apache.hadoop.io.IntWritable;    import
org.apache.hadoop.io.Text;    import    org.apache.hadoop.mapreduce.Job;    import
org.apache.hadoop.mapreduce.lib.input.FileInputFormat;    import
org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```
public class AverageDriver {
    public static void main(String[] args) throws Exception {
if (args.length != 2) {
        System.err.println("Please Enter the input and output parameters");
        System.exit(-1);
    }
    Job job = new Job(); job.setJarByClass(AverageDriver.class);
    job.setJobName("Max temperature");
    FileInputFormat.addInputPath(job,    new    Path(args[0]));
    FileOutputFormat.setOutputPath(job,    new    Path(args[1]));
    job.setMapperClass(AverageMapper.class); job.setReducerClass(AverageReducer.class);
    job.setOutputKeyClass(Text.class);    job.setOutputValueClass(IntWritable.class);
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

## AverageMapper

**package** temp;

```
import java.io.IOException; import org.apache.hadoop.io.IntWritable;
import    org.apache.hadoop.io.LongWritable;    import
org.apache.hadoop.io.Text;    import
org.apache.hadoop.mapreduce.Mapper; public class AverageMapper
extends
```

```

Mapper<LongWritable, Text, Text, IntWritable> {    public
static final int MISSING = 9999;

    public void map(LongWritable key, Text value, Mapper<LongWritable, Text, Text,
IntWritable>.Context context) throws IOException, InterruptedException {    int
temperature;
    String line = value.toString();    String year = line.substring(15, 19);    if
(line.charAt(87) == '+') { temperature = Integer.parseInt(line.substring(88, 92));
    } else {
        temperature = Integer.parseInt(line.substring(87, 92));
    }
    String quality = line.substring(92, 93);
    if (temperature != 9999 && quality.matches("[01459]")) context.write(new
        Text(year), new
        IntWritable(temperature));
    }
}

```

AverageReducer **package**

temp;

```

import    java.io.IOException;    import
org.apache.hadoop.io.IntWritable;    import
org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

```

```

public class AverageReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
public void reduce(Text key, Iterable<IntWritable> values, Reducer<Text, IntWritable, Text,
IntWritable>.Context context) throws IOException, InterruptedException {    int max_temp = 0;
    int count = 0;

```

```
    for (IntWritable value : values)
    {
        max_temp += value.get(); count++;
    }
    context.write(key, new IntWritable(max_temp / count));
}
}
```

Output:

```
hduser@bmsce-Precision-T1700:~/Desktop/temperature$ start-all.sh
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
Starting namenodes on [localhost]
hduser@localhost's password:
localhost: starting namenode, logging to /usr/local/hadoop/logs/hadoop-hduser-namenode-bmsce-
Precision-T1700.out
hduser@localhost's password:
localhost: starting datanode, logging to /usr/local/hadoop/logs/hadoop-hduser-datanode-bmsce-
Precision-T1700.out
Starting secondary namenodes [0.0.0.0]
hduser@0.0.0.0's password:
0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop/logs/hadoop-hduser-
secondarynamenode-bmsce-Precision-T1700.out
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop/logs/yarn-hduser-resourcemanager-bmsce-
Precision-T1700.out
hduser@localhost's password:
localhost: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-hduser-nodemanager-bmsce-
Precision-T1700.out
hduser@bmsce-Precision-T1700:~/Desktop/temperature$ jps
6832 NodeManager
6498 ResourceManager
6339 SecondaryNameNode
4887 org.eclipse.equinox.launcher_1.5.600.v20191014-2022.jar
6954 Jps
6123 DataNode
5951 NameNode
hduser@bmsce-Precision-T1700:~/Desktop/temperature$ hdfs dfs -ls /
-ls: Unknown command
hduser@bmsce-Precision-T1700:~/Desktop/temperature$ hdfs dfs -ls /
Found 31 items
drwxr-xr-x - hduser supergroup 0 2022-06-06 12:35 /CSE
drwxr-xr-x - hduser supergroup 0 2022-06-06 12:23 /FFF
drwxr-xr-x - hduser supergroup 0 2022-06-06 12:36 /LLL
drwxr-xr-x - hduser supergroup 0 2022-06-20 12:06 /amit_bda
drwxr-xr-x - hduser supergroup 0 2022-06-27 11:42 /amit_lab
drwxr-xr-x - hduser supergroup 0 2022-06-03 14:52 /bharath
drwxr-xr-x - hduser supergroup 0 2022-06-03 14:43 /bharath035
drwxr-xr-x - hduser supergroup 0 2022-06-24 14:54 /chi
drwxr-xr-x - hduser supergroup 0 2022-05-31 10:21 /example
drwxr-xr-x - hduser supergroup 0 2022-06-01 15:13 /foldernew
drwxr-xr-x - hduser supergroup 0 2022-06-06 15:04 /hemang061
drwxr-xr-x - hduser supergroup 0 2022-06-20 15:16 /input_khushil
drwxr-xr-x - hduser supergroup 0 2022-06-03 12:27 /irfan
drwxr-xr-x - hduser supergroup 0 2022-06-22 10:44 /lwde
drwxr-xr-x - hduser supergroup 0 2022-06-27 13:03 /mapreducejoin_amit
drwxr-xr-x - hduser supergroup 0 2022-06-22 15:32 /muskan
drwxr-xr-x - hduser supergroup 0 2022-06-22 15:06 /muskan_op
drwxr-xr-x - hduser supergroup 0 2022-06-22 15:35 /muskan_output
drwxr-xr-x - hduser supergroup 0 2022-06-06 15:04 /new_folder
drwxr-xr-x - hduser supergroup 0 2022-05-31 10:26 /one
drwxr-xr-x - hduser supergroup 0 2022-06-24 15:30 /out55
drwxr-xr-x - hduser supergroup 0 2022-06-20 12:17 /output
drwxr-xr-x - hduser supergroup 0 2022-06-27 13:04 /output_TOPn
drwxr-xr-x - hduser supergroup 0 2022-06-27 12:14 /output_Topn
drwxr-xr-x - hduser supergroup 0 2022-06-24 12:42 /r1
drwxr-xr-x - hduser supergroup 0 2022-06-24 12:24 /rgs
```